



held-out text loss

lower is better



held-out image loss

lower is better



the cross-modal test

loss on tokens only the other modality explains

