



What multimodal training costs the language model underneath



Same images, same captions, same cross-modal dependency in every arm. The only difference is how much of the text is prose rather than a caption.